

problem Statement: To predict How Best the data fits

1) Data Collection

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
```

```
In [2]: df = pd.read_csv(r"C:\Users\chinta pavani\Documents\insurance.csv")
df
```

Out[2]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

2) Data Cleaning and Preprocessing

In [3]: `df.head()`

Out[3]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

In [4]: `df.tail()`

Out[4]:

	age	sex	bmi	children	smoker	region	charges
1333	50	male	30.97	3	no	northwest	10600.5483
1334	18	female	31.92	0	no	northeast	2205.9808
1335	18	female	36.85	0	no	southeast	1629.8335
1336	21	female	25.80	0	no	southwest	2007.9450
1337	61	female	29.07	0	yes	northwest	29141.3603

In [5]: `df.shape`

Out[5]: (1338, 7)

In [6]: `df.describe`

Out[6]:

```
<bound method NDFrame.describe of
ker      region      charges
0      19  female  27.900      0  yes  southwest  16884.92400
1      18   male  33.770      1  no   southeast  1725.55230
2      28   male  33.000      3  no   southeast  4449.46200
3      33   male  22.705      0  no   northwest  21984.47061
4      32   male  28.880      0  no   northwest  3866.85520
...     ...     ...      ...     ...     ...     ...
1333    50   male  30.970      3  no   northwest  10600.54830
1334    18  female  31.920      0  no   northeast  2205.98080
1335    18  female  36.850      0  no   southeast  1629.83350
1336    21  female  25.800      0  no   southwest  2007.94500
1337    61  female  29.070      0  yes  northwest  29141.36030

[1338 rows x 7 columns]>
```

In [7]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype  
---  -
 0   age         1338 non-null   int64  
 1   sex         1338 non-null   object  
 2   bmi         1338 non-null   float64  
 3   children    1338 non-null   int64  
 4   smoker      1338 non-null   object  
 5   region      1338 non-null   object  
 6   charges     1338 non-null   float64  
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

In [8]: `df.isnull().any()`

```
Out[8]: age         False
sex         False
bmi         False
children    False
smoker      False
region      False
charges     False
dtype: bool
```

In [9]: `df.isna().sum()`

```
Out[9]: age         0
sex         0
bmi         0
children    0
smoker      0
region      0
charges     0
dtype: int64
```

In [10]: `df['region'].value_counts()`

```
Out[10]: region
southeast    364
southwest    325
northwest    325
northeast    324
Name: count, dtype: int64
```

```
In [11]: ▶ convert={"sex":{"female":1,"male":0}}
df=df.replace(convert)
df
```

Out[11]:

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	yes	southwest	16884.92400
1	18	0	33.770	1	no	southeast	1725.55230
2	28	0	33.000	3	no	southeast	4449.46200
3	33	0	22.705	0	no	northwest	21984.47061
4	32	0	28.880	0	no	northwest	3866.85520
...
1333	50	0	30.970	3	no	northwest	10600.54830
1334	18	1	31.920	0	no	northeast	2205.98080
1335	18	1	36.850	0	no	southeast	1629.83350
1336	21	1	25.800	0	no	southwest	2007.94500
1337	61	1	29.070	0	yes	northwest	29141.36030

```
In [12]: ▶ convert={"smoker":{"yes":1, "no":0}}
df=df.replace(convert)
df
```

Out[12]:

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	1	southwest	16884.92400
1	18	0	33.770	1	0	southeast	1725.55230
2	28	0	33.000	3	0	southeast	4449.46200
3	33	0	22.705	0	0	northwest	21984.47061
4	32	0	28.880	0	0	northwest	3866.85520
...
1333	50	0	30.970	3	0	northwest	10600.54830
1334	18	1	31.920	0	0	northeast	2205.98080
1335	18	1	36.850	0	0	southeast	1629.83350
1336	21	1	25.800	0	0	southwest	2007.94500
1337	61	1	29.070	0	1	northwest	29141.36030

1338 rows × 7 columns

```
In [13]: ► convert={"region":{"southwest":1, "southeast":2,"northwest":3,"northeast":4}
df=df.replace(convert)
df
```


Out[13]:

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	1	1	16884.92400
1	18	0	33.770	1	0	2	1725.55230
2	28	0	33.000	3	0	2	4449.46200
3	33	0	22.705	0	0	3	21984.47061
4	32	0	28.880	0	0	3	3866.85520
...
1333	50	0	30.970	3	0	3	10600.54830
1334	18	1	31.920	0	0	4	2205.98080
1335	18	1	36.850	0	0	2	1629.83350
1336	21	1	25.800	0	0	1	2007.94500
1337	61	1	29.070	0	1	3	29141.36030

1338 rows × 7 columns

3)Data Visualization

```
In [14]: sns.lmplot(x='bmi',y='charges',order=2,data=df,ci=None)
plt.show()
```




```
In [15]: x=np.array(df['bmi']).reshape(-1,1)
y=np.array(df['charges']).reshape(-1,1)
```

```
In [16]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25,random_s
lr=LinearRegression()
lr.fit(x_train,y_train)
print(lr.score(x_test,y_test))


0.050136213258239914
```

```
In [17]: ▶ y_pred=lr.predict(x_test)
plt.scatter(x_test,y_test,color='b')
plt.plot(x_test,y_pred,color='k')
plt.show()
```



Working With Subset Of Data

```
In [18]: ▶ df700=df[:][:700]
sns.lmplot(x="bmi",y="charges",order=2,ci=None,data=df700)
plt.show()
```



```
In [19]: ▶ df700.fillna(method='ffill',inplace=True)
```


```
In [20]: x=np.array(df['bmi']).reshape(-1,1)
         y=np.array(df['charges']).reshape(-1,1)
```

```
In [21]: df700.dropna(inplace=True)
```

```
In [22]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
         lr=LinearRegression()
         lr.fit(x_train,y_train)
         print(lr.score(x_test,y_test))

0.03834819458441918
```

```
In [23]: y_pred=lr.predict(x_test)
         plt.scatter(x_test,y_test,color='b')
         plt.plot(x_test,y_pred,color='k')
         plt.show()
```



Evaluation of model

```
In [24]: from sklearn.linear_model import LinearRegression
         from sklearn.metrics import r2_score
         lr=LinearRegression()
         lr.fit(x_train,y_train)
         y_pred=lr.predict(x_test)
         r2=r2_score(y_test,y_pred)
         print(r2)
```

0.03834819458441918

Ridge Regression


```
In [25]: from sklearn.linear_model import Lasso,Ridge
from sklearn.preprocessing import StandardScaler
```

```
In [26]: plt.figure(figsize= (10,10))
sns.heatmap(df700.corr(),annot = True)
plt.show()
```



```
In [27]: features =df.columns[0:1]
target = df.columns[-1]
x = df[features].values
y = df[target].values
x_train,x_test,y_train,y_test = train_test_split(x,y, test_size=0.30,random
print("The dimension of x_train is {}".format(x_train.shape))
print("The dimension of x_test is {}".format(x_test.shape))
```

```
The dimension of x_train is (936, 1)
The dimension of x_test is (402, 1)
```

```
In [28]: ▶ lr = LinearRegression()
#Fit model
lr.fit(x_train,y_train)
#Predict
#Prediction =Lr.predict(x_test)
#actual
actual = y_test
train_score_lr = lr.score(x_train,y_train)
test_score_lr = lr.score(x_test,y_test)
print("\nLinear Regression Model :\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model :


The train score for lr model is 0.0910963973805714
The test score for lr model is 0.08490473916580776

```
In [29]: ▶ ridgeReg = Ridge(alpha=10)
ridgeReg.fit(x_train,y_train)
#train and test score for ridge regression
train_score_ridge = ridgeReg.score(x_train,y_train)
test_score_ridge = ridgeReg.score(x_test,y_test)
print("\nRidge Model :\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model :

The train score for ridge model is 0.09109639711159634
The test score for ridge model is 0.08490538609860176

```
In [30]: ▶ plt.figure(figsize =(10,10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',mar
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize
plt.xticks(rotation =90)
plt.legend()
plt.show()
```



Lasso Regression


```
In [31]: ▶ lasso=Lasso(alpha=10)
lasso.fit(x_train,y_train)
train_score_ls =lasso.score(x_train,y_train)
test_score_ls =lasso.score(x_test,y_test)
print("\nRidge Model:\n")
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))
```

Ridge Model:

The train score for ls model is 0.09109639395809055

The test score for ls model is 0.08490704421828055

```
In [32]: ▶ plt.figure(figsize=(10,10))
pd.Series(lasso.coef_, features).sort_values(ascending =True).plot(kind =
plt.show())
```



```
In [33]: ▶ from sklearn.linear_model import LassoCV
```


```
In [34]: #using the Linear CV model
from sklearn.linear_model import RidgeCV
#Lasso Cross validation
ridge_cv = RidgeCV(alphas = [0.0001,0.001,0.01,0.1,1,10]).fit(x_train,y_train)
#score
print(ridge_cv.score(x_train,y_train))
print(ridge_cv.score(x_test,y_test))
```

```
0.09109639711159612
0.08490538609884779
```

```
In [47]: #using the linear CV model
from sklearn.linear_model import LassoCV
#Ridge Cross validation
lasso_cv=LassoCV(alphas = [0.0001,0.001,0.01,1,10]).fit(x_train,y_train)
#score
print("The train score for ls model is {}".format(ridge_cv.score(x_train,y_train)))
print("The test score for ls model is {}".format(ridge_cv.score(x_test,y_test)))
```

```
The train score for ls model is 0.09109639711159612
The test score for ls model is 0.08490538609884779
```

```
In [48]: plt.figure(figsize = (10, 10))
#add plot for ridge regression
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=6)
#add plot for lasso regression
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6)
#add plot for linear model
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=6)
#rotate axis
plt.xticks(rotation = 90)
plt.legend()
plt.title("Comparison plot of Ridge, Lasso and Linear regression model")
plt.show()
```



Elastic Net Regression

```
In [49]: ▶ from sklearn.linear_model import ElasticNet
          el=ElasticNet()
          el.fit(x,y)
          print(el.coef_)
          print(el.intercept_)
```

```
[257.0684655]
3191.532406056678
```

```
In [50]: ▶ y_pred_elastic=el.predict(x_train)
```

```
In [51]: ▶ mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
          print("Mean Squared Error on test",mean_squared_error)
```

```
Mean Squared Error on test 135093713.48214507
```

```
In [52]: ▶ el=ElasticNet()
          el.fit(x_train,y_train)
          print(el.score(x_train,y_train))
```

```
0.09109580670592365
```

Logistic Regression

```
In [80]: ▶ import numpy as np
          import pandas as pd
          from sklearn.linear_model import LogisticRegression
          from sklearn.preprocessing import StandardScaler
```

```
In [81]: df=pd.read_csv(r"C:\Users\chinta pavani\Documents\insurance.csv")
df
```

Out[81]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.924000
1	18	male	33.770	1	no	southeast	1725.552300
2	28	male	33.000	3	no	southeast	4449.462000
3	33	male	22.705	0	no	northwest	21984.470610
4	32	male	28.880	0	no	northwest	3866.855200
5	31	female	25.740	0	no	southeast	3756.621600
6	46	female	33.440	1	no	southeast	8240.589600
7	37	female	27.740	3	no	northwest	7281.505600
8	37	male	29.830	2	no	northeast	6406.410700
9	60	female	25.840	0	no	northwest	28923.136920
10	25	male	26.220	0	no	northeast	2721.320800

```
In [82]: df.shape
```

Out[82]: (1338, 7)

```
In [83]: pd.set_option('display.max_rows',10000000000)
pd.set_option('display.max_columns',10000000000)
pd.set_option('display.width',95)
```

```
In [84]: print('This Dataset has %d rows and %d columns'%(df.shape))
```

This Dataset has 1338 rows and 7 columns

```
In [85]: df.head()
```

Out[85]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

In [86]: `df.describe`

Out[86]:

	ker	region	charges		age	sex	bmi	children	smo
0	19	female	27.900	0	yes	southwest	16884.924000		
1	18	male	33.770	1	no	southeast	1725.552300		
2	28	male	33.000	3	no	southeast	4449.462000		
3	33	male	22.705	0	no	northwest	21984.470610		
4	32	male	28.880	0	no	northwest	3866.855200		
5	31	female	25.740	0	no	southeast	3756.621600		
6	46	female	33.440	1	no	southeast	8240.589600		
7	37	female	27.740	3	no	northwest	7281.505600		
8	37	male	29.830	2	no	northeast	6406.410700		
9	60	female	25.840	0	no	northwest	28923.136920		
10	25	male	26.220	0	no	northeast	2721.320800		
11	62	female	26.290	0	yes	southeast	27808.725100		
12	23	male	34.400	0	no	southwest	1826.843000		
13	56	female	39.820	0	no	southeast	11090.717800		
14	27	male	42.130	0	yes	southeast	39611.757700		
15	19	male	24.600	1	no	southwest	1837.237000		
16	52	female	30.780	1	no	northeast	10797.336200		
17	22	male	22.045	0	no	southwest	2205.471550		

In [87]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

In [88]: `df.isnull().sum()`

Out[88]:

```
age         0
sex         0
bmi         0
children    0
smoker      0
region      0
charges     0
dtype: int64
```



```
In [89]: ▶ convert={"smoker":{"yes":1,"no":0}}
df=df.replace(convert)
df
```

Out[89]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	1	southwest	16884.924000
1	18	male	33.770	1	0	southeast	1725.552300
2	28	male	33.000	3	0	southeast	4449.462000
3	33	male	22.705	0	0	northwest	21984.470610
4	32	male	28.880	0	0	northwest	3866.855200
5	31	female	25.740	0	0	southeast	3756.621600
6	46	female	33.440	1	0	southeast	8240.589600
7	37	female	27.740	3	0	northwest	7281.505600
8	37	male	29.830	2	0	northeast	6406.410700
9	60	female	25.840	0	0	northwest	28923.136920
10	25	male	26.220	0	0	northeast	2721.320800

```
In [90]: ▶ convert={"sex":{"female":1,"male":0}}
df=df.replace(convert)
df
```

Out[90]:

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	1	southwest	16884.924000
1	18	0	33.770	1	0	southeast	1725.552300
2	28	0	33.000	3	0	southeast	4449.462000
3	33	0	22.705	0	0	northwest	21984.470610
4	32	0	28.880	0	0	northwest	3866.855200
5	31	1	25.740	0	0	southeast	3756.621600
6	46	1	33.440	1	0	southeast	8240.589600
7	37	1	27.740	3	0	northwest	7281.505600
8	37	0	29.830	2	0	northeast	6406.410700
9	60	1	25.840	0	0	northwest	28923.136920
10	25	0	26.220	0	0	northeast	2721.320800

```
In [91]: ▶ convert={"region":{"southeast":1,"southwest":2,"northeast":3,"northwest":4}
df=df.replace(convert)
df
```

Out[91]:

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	1	2	16884.924000
1	18	0	33.770	1	0	1	1725.552300
2	28	0	33.000	3	0	1	4449.462000
3	33	0	22.705	0	0	4	21984.470610
4	32	0	28.880	0	0	4	3866.855200
5	31	1	25.740	0	0	1	3756.621600
6	46	1	33.440	1	0	1	8240.589600
7	37	1	27.740	3	0	4	7281.505600
8	37	0	29.830	2	0	3	6406.410700
9	60	1	25.840	0	0	4	28923.136920
10	25	0	26.220	0	0	3	2721.320800

```
In [92]: ▶ features_matrix=df.iloc[:,0:4]
```


```
In [93]: ▶ target_vector=df.iloc[:, -3]
```

```
In [94]: ▶ print('The Feature Matrix has %d Rows and %d columns(s)'%(features_matrix.
print('The Target Matrix has %d Rows and %d columns(s)'%(np.array(target_v
```

The Feature Matrix has 1338 Rows and 4 columns(s)

The Target Matrix has 1338 Rows and 1 columns(s)

```
In [95]: import matplotlib.pyplot as plt
import seaborn as sns
sns.barplot(x='smoker', y='age', data=df, color="mediumturquoise")
plt.show()
```



```
In [136]: features_matrix_standardized=StandardScaler().fit_transform(features_matrix)
```

```
In [137]: algorithm=LogisticRegression(max_iter=10000)
```

```
In [138]: Logistic_Regression_Model=algorithm.fit(features_matrix_standardized,target)
```

```
In [139]: observation=[[1,0,0.99539,-0.0588]]
```

```
In [140]: predictions=Logistic_Regression_Model.predict(observation)
print('The model predicted the observation to belong to class %s'%(predictions[0]))
The model predicted the observation to belong to class [0]
```

```
In [141]: print('The algorithm was trained to predict one of the two classes:%s'%(algorithm.classes_))
The algorithm was trained to predict one of the two classes:[0 1]
```

```
In [142]: ▶ print(" " "The Model says the probability of the observation we passed belong to class[0] 0.8057075871331396")
print()
```

The Model says the probability of the observation we passed belonging to class[0] 0.8057075871331396

```
In [148]: ▶ print(" " "The Model says the probability of the observation we passed belong to class['g'] is 0.19429241286686041")
print()
```

The Model says the probability of the observation we passed belonging to class['g'] is 0.19429241286686041

```
In [146]: ▶ x=np.array(df['age']).reshape(-1,1)
y=np.array(df['smoker']).reshape(-1,1)
```

```
In [102]: ▶ x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.05)
lo=LogisticRegression()
lo.fit(x_train,y_train)
print(lo.score(x_test,y_test))
```

0.7910447761194029

C:\Users\chinta pavani\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\utils\validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().
y = column_or_1d(y, warn=True)

Decision Tree

```
In [103]: ▶ import numpy as np
import pandas as pd
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
```

```
In [104]: df=pd.read_csv(r"C:\Users\chinta pavani\Documents\insurance.csv")
df
```

Out[104]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.924000
1	18	male	33.770	1	no	southeast	1725.552300
2	28	male	33.000	3	no	southeast	4449.462000
3	33	male	22.705	0	no	northwest	21984.470610
4	32	male	28.880	0	no	northwest	3866.855200
5	31	female	25.740	0	no	southeast	3756.621600
6	46	female	33.440	1	no	southeast	8240.589600
7	37	female	27.740	3	no	northwest	7281.505600
8	37	male	29.830	2	no	northeast	6406.410700
9	60	female	25.840	0	no	northwest	28923.136920
10	25	male	26.220	0	no	northeast	2721.320800

```
In [105]: df.shape
```

Out[105]: (1338, 7)

```
In [106]: df.isnull().any()
```

Out[106]:

age	False
sex	False
bmi	False
children	False
smoker	False
region	False
charges	False

dtype: bool

```
In [107]: df['region'].value_counts()
```

Out[107]:

region	
southeast	364
southwest	325
northwest	325
northeast	324

Name: count, dtype: int64

```
In [108]: ▶ convert={"sex":{"female":1,"male":0}}
df=df.replace(convert)
df
```

Out[108]:

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	yes	southwest	16884.924000
1	18	0	33.770	1	no	southeast	1725.552300
2	28	0	33.000	3	no	southeast	4449.462000
3	33	0	22.705	0	no	northwest	21984.470610
4	32	0	28.880	0	no	northwest	3866.855200
5	31	1	25.740	0	no	southeast	3756.621600
6	46	1	33.440	1	no	southeast	8240.589600
7	37	1	27.740	3	no	northwest	7281.505600
8	37	0	29.830	2	no	northeast	6406.410700
9	60	1	25.840	0	no	northwest	28923.136920
10	25	0	26.220	0	no	northeast	2721.320800

```
In [109]: ▶ convert={"smoker":{"yes":1,"no":0}}
df=df.replace(convert)
df
```

Out[109]:

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	1	southwest	16884.924000
1	18	0	33.770	1	0	southeast	1725.552300
2	28	0	33.000	3	0	southeast	4449.462000
3	33	0	22.705	0	0	northwest	21984.470610
4	32	0	28.880	0	0	northwest	3866.855200
5	31	1	25.740	0	0	southeast	3756.621600
6	46	1	33.440	1	0	southeast	8240.589600
7	37	1	27.740	3	0	northwest	7281.505600
8	37	0	29.830	2	0	northeast	6406.410700
9	60	1	25.840	0	0	northwest	28923.136920
10	25	0	26.220	0	0	northeast	2721.320800

```
In [110]: ▶ x=["bmi","children"]
y=["Yes","No"]
all_inputs=df[x]
all_classes=df["sex"]
```

```
In [111]: ▶ (x_train,x_test,y_train,y_test)=train_test_split(all_inputs,all_classes,te
```

```
In [112]: clf=DecisionTreeClassifier(random_state=0)
```

```
In [113]: clf.fit(x_train,y_train)
```

```
Out[113]: DecisionTreeClassifier
DecisionTreeClassifier(random_state=0)
```

```
In [114]: DecisionTreeClassifier(random_state=0)
score=clf.score(x_test,y_test)
print(score)
```

```
0.4878048780487805
```

Random Forest

```
In [115]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt ,seaborn as sns
```

```
In [116]: df=pd.read_csv(r"C:\Users\chinta pavani\Documents\insurance.csv")
df
```

```
Out[116]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.924000
1	18	male	33.770	1	no	southeast	1725.552300
2	28	male	33.000	3	no	southeast	4449.462000
3	33	male	22.705	0	no	northwest	21984.470610
4	32	male	28.880	0	no	northwest	3866.855200
5	31	female	25.740	0	no	southeast	3756.621600
6	46	female	33.440	1	no	southeast	8240.589600
7	37	female	27.740	3	no	northwest	7281.505600
8	37	male	29.830	2	no	northeast	6406.410700
9	60	female	25.840	0	no	northwest	28923.136920
10	25	male	26.220	0	no	northeast	2721.320800

```
In [117]: df.shape
```

```
Out[117]: (1338, 7)
```

```
In [118]: df['region'].value_counts()
```

```
Out[118]: region
southeast    364
southwest    325
northwest    325
northeast    324
Name: count, dtype: int64
```

```
In [119]: df['bmi'].value_counts()
```

```
Out[119]: bmi
32.300    13
28.310     9
30.495     8
30.875     8
31.350     8
30.800     8
34.100     8
28.880     8
33.330     7
35.200     7
25.800     7
32.775     7
27.645     7
32.110     7
38.060     7
25.460     7
30.590     7
27.360     7
24.330     7
```

```
In [121]: m={"sex":{"female":1,"male":0}}
df=df.replace(m)
print(df)
```

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	yes	southwest	16884.924000
1	18	0	33.770	1	no	southeast	1725.552300
2	28	0	33.000	3	no	southeast	4449.462000
3	33	0	22.705	0	no	northwest	21984.470610
4	32	0	28.880	0	no	northwest	3866.855200
5	31	1	25.740	0	no	southeast	3756.621600
6	46	1	33.440	1	no	southeast	8240.589600
7	37	1	27.740	3	no	northwest	7281.505600
8	37	0	29.830	2	no	northeast	6406.410700
9	60	1	25.840	0	no	northwest	28923.136920
10	25	0	26.220	0	no	northeast	2721.320800
11	62	1	26.290	0	yes	southeast	27808.725100
12	23	0	34.400	0	no	southwest	1826.843000
13	56	1	39.820	0	no	southeast	11090.717800
14	27	0	42.130	0	yes	southeast	39611.757700
15	19	0	24.600	1	no	southwest	1837.237000
16	52	1	30.780	1	no	northeast	10797.336200
17	23	0	23.845	0	no	northeast	2395.171550
18	55	0	40.300	0	no	southwest	10603.265000


```
In [122]: n={"smoker":{"yes":1,"no":0}}
df=df.replace(n)
print(df)
```

	age	sex	bmi	children	smoker	region	charges
0	19	1	27.900	0	1	southwest	16884.924000
1	18	0	33.770	1	0	southeast	1725.552300
2	28	0	33.000	3	0	southeast	4449.462000
3	33	0	22.705	0	0	northwest	21984.470610
4	32	0	28.880	0	0	northwest	3866.855200
5	31	1	25.740	0	0	southeast	3756.621600
6	46	1	33.440	1	0	southeast	8240.589600
7	37	1	27.740	3	0	northwest	7281.505600
8	37	0	29.830	2	0	northeast	6406.410700
9	60	1	25.840	0	0	northwest	28923.136920
10	25	0	26.220	0	0	northeast	2721.320800
11	62	1	26.290	0	1	southeast	27808.725100
12	23	0	34.400	0	0	southwest	1826.843000
13	56	1	39.820	0	0	southeast	11090.717800
14	27	0	42.130	0	1	southeast	39611.757700
15	19	0	24.600	1	0	southwest	1837.237000
16	52	1	30.780	1	0	northeast	10797.336200
17	23	0	23.845	0	0	northeast	2395.171550
18	55	0	40.200	0	0	southwest	10603.265000

```
In [123]: from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[123]: ▼ RandomForestClassifier
RandomForestClassifier()
```

```
In [124]: rf=RandomForestClassifier()
params={'max_depth':[2,3,5,20],
'min_samples_leaf':[5,10,20,50,100,200],
'n_estimators':[10,25,30,50,100,200]}
```

```
In [126]: from sklearn.model_selection import GridSearchCV
grid_search=GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring="accu
grid_search.fit(x_train,y_train)
```

```
Out[126]: ▶ GridSearchCV
▶ estimator: RandomForestClassifier
▶ RandomForestClassifier
```

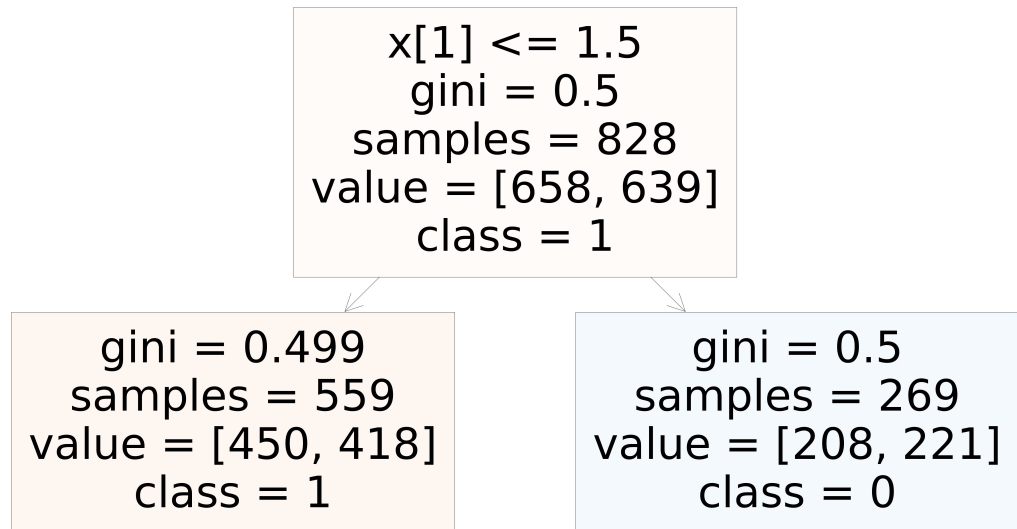
```
In [127]: grid_search.best_score_
```

```
Out[127]: 0.5304492666780802
```

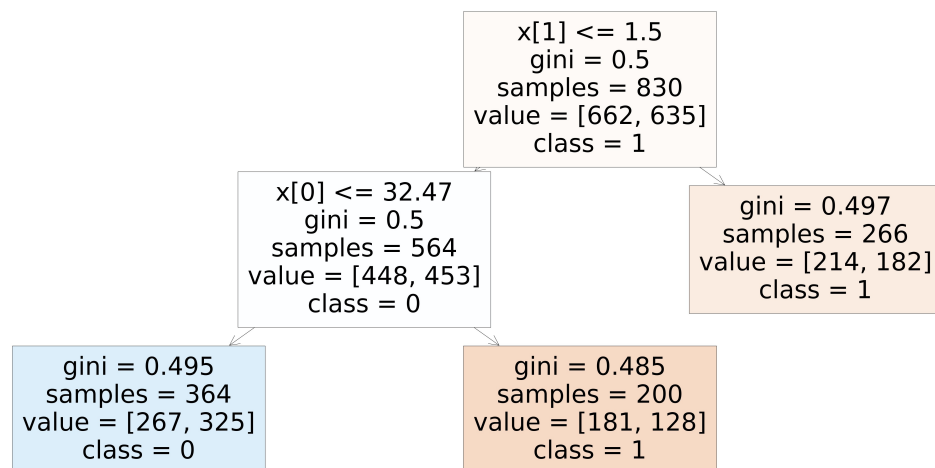
```
In [128]: rf_best=grid_search.best_estimator_
print(rf_best)
```

```
RandomForestClassifier(max_depth=3, min_samples_leaf=200, n_estimators=10)
```

```
In [129]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[4],class_names=['1','0'],filled=True);
```




```
In [131]: from sklearn.tree import plot_tree
plt.figure(figsize=(70,30))
plot_tree(rf_best.estimators_[6],class_names=["1","0"],filled=True);
```



```
In [132]: rf_best.feature_importances_
```


```
Out[132]: array([0.53541719, 0.46458281])
```

```
In [133]: rf=RandomForestClassifier(random_state=0)
```

In [134]:  rf.fit(x_train,y_train)

Out[134]:

▼ RandomForestClassifier
RandomForestClassifier(random_state=0)

In [135]:  score=rf.score(x_test,y_test)
print(score)

0.4878048780487805

In []: 